




UNIVERSITY OF COPENHAGEN

Communication: Message Queues & BASE

ACS, Yongluan Zhou


What should we learn today?



- Describe different approaches to design **communication abstractions**, e.g., transient vs. persistent and synchronous vs. asynchronous
- Explain the design and implementation of **message-oriented middleware** (MOM)
- Explain how to organize systems employing a **BASE** methodology
- Discuss the relationship of BASE to eventual consistency and the **CAP theorem**
- Identify alternative communication abstractions such as data **streams and multicast** / gossip


Towards Distribution

- Use independent services to store different data → **partitioning**
- Scalability and Availability improves but now coordination may be necessary
- Employ communication abstraction




Towards Distribution

- Use independent services to store different data → **partitioning**
- Scalability and Availability improves but now coordination may be necessary
- Employ communication abstraction



Recall: Fundamental abstractions


- Memory
 - Read/write



- Interpreter
 - Instruction repertoire
 - Environment
 - Instruction pointer

`(loop (print (eval (read))))`

Networks
implement link
abstraction



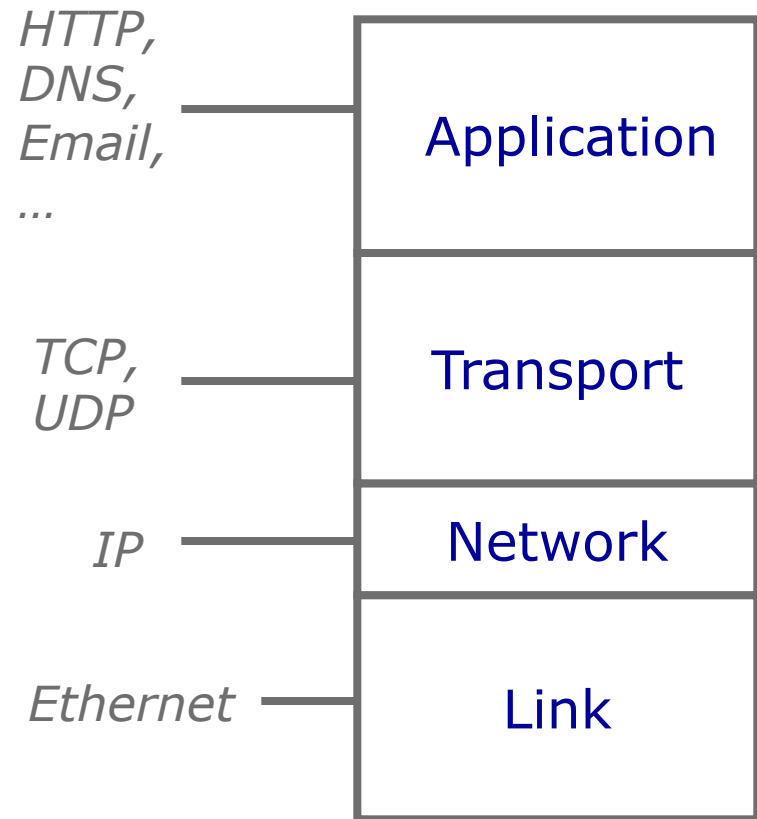
- Communication links
 - Send/receive

Source: Saltzer & Kaashoek & Morris (partial)



Recall: Protocols and Layering in the Internet

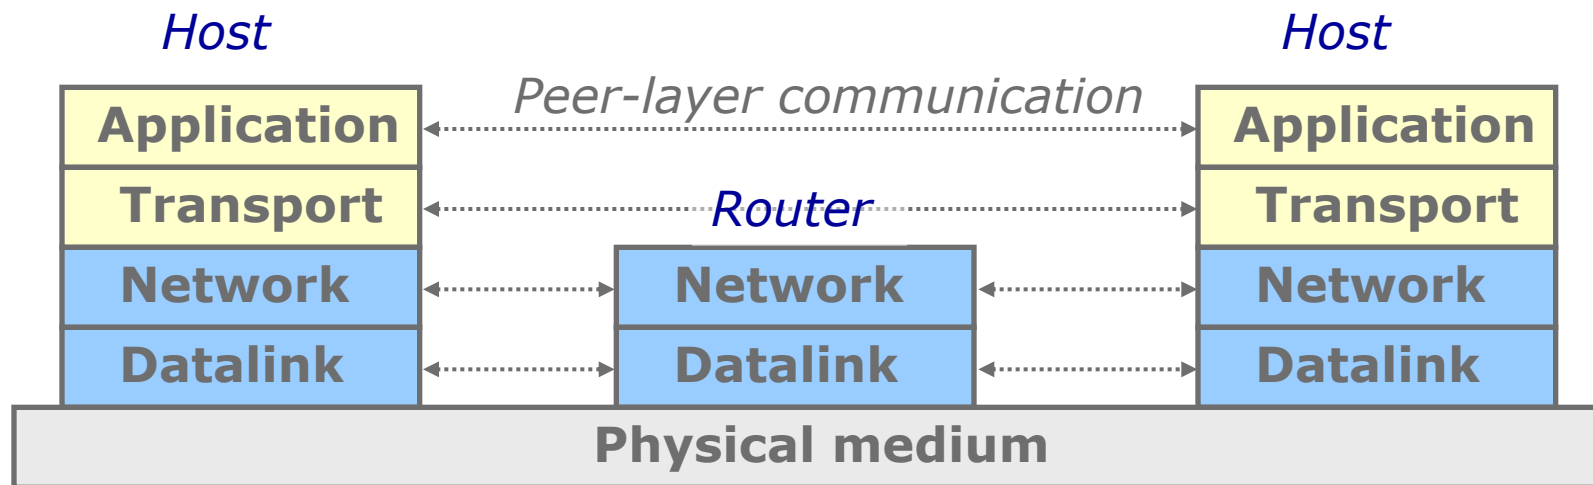
- Layering
 - System broken into vertical hierarchy of protocols
 - Service provided by one layer based solely on service provided by layer below
- Internet model
 - Four layers



Do all servers in the network contain all layers?

Recall: Layers in Hosts and Routers

- Link and network layers implemented everywhere
- End-to-end layer (i.e., transport and application) implemented only at hosts



Source: Tanenbaum &
Van Steen

Different Types of Communication

- Synchronous vs. Asynchronous
 - Synchronize at request submission (1)
 - Synchronize at request delivery (2)
 - Synchronize after being fully processed by recipient (3)
- Transient vs. Persistent (temporal decoupling)

Examples	Persistent	Transient
Asynchronous	Email	UDP Erlang
Synchronous	Message-queueing systems (1) Online Chat (1) + (2)	Asynchronous RPC (2) RPC (3)



Two Types of Communication Abstractions

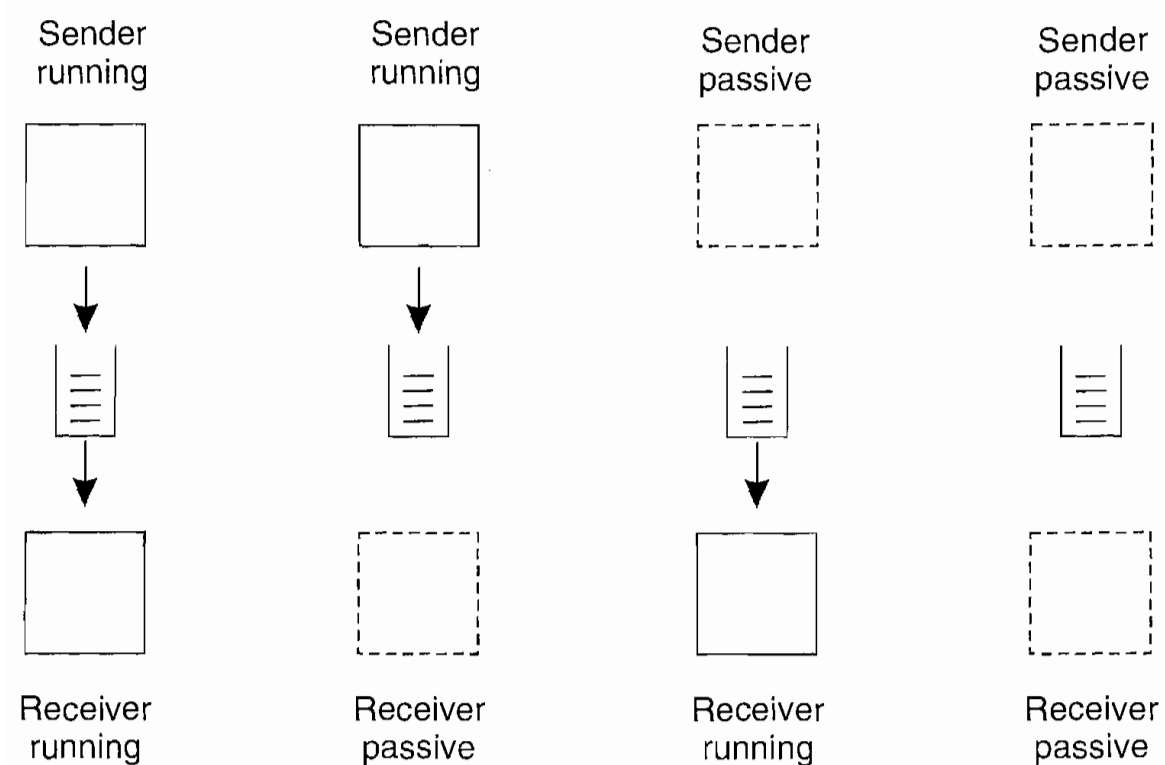
- **RPC**
 - Transparent and hidden communication
 - Synchronous
 - Transient
- **Message-Oriented**
 - Explicit communication SEND/RECEIVE of point-to-point messages
 - Synchronous vs. Asynchronous
 - Transient vs. Persistent



Source: Saltzer & Kaashoek & Morris (partial),
Tanenbaum & Steen (partial)

Message-Oriented Persistent Communication

- Queues make sender and receiver loosely-coupled
- Modes of execution of sender/receiver:

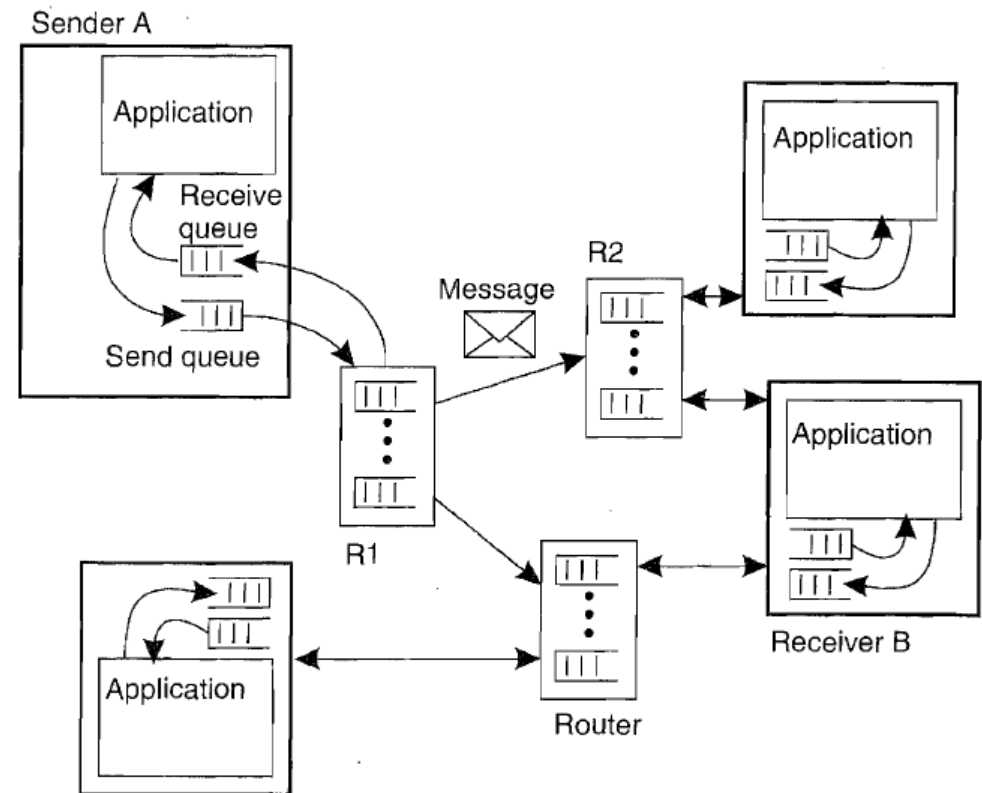


Source: Tanenbaum & Van Steen



Queue Interface

Put	Put message in queue
Get	Remove first message from queue (blocking)
Poll	Check for message and remove first (non-blocking)
Notify	Handler that is called when message is added



- Source / destination decoupled by **queue names**
- **Relays:** store and forward messages
- **Brokers:** gateway to transform message formats

Questions so far?

Employing Queues to Decouple Systems

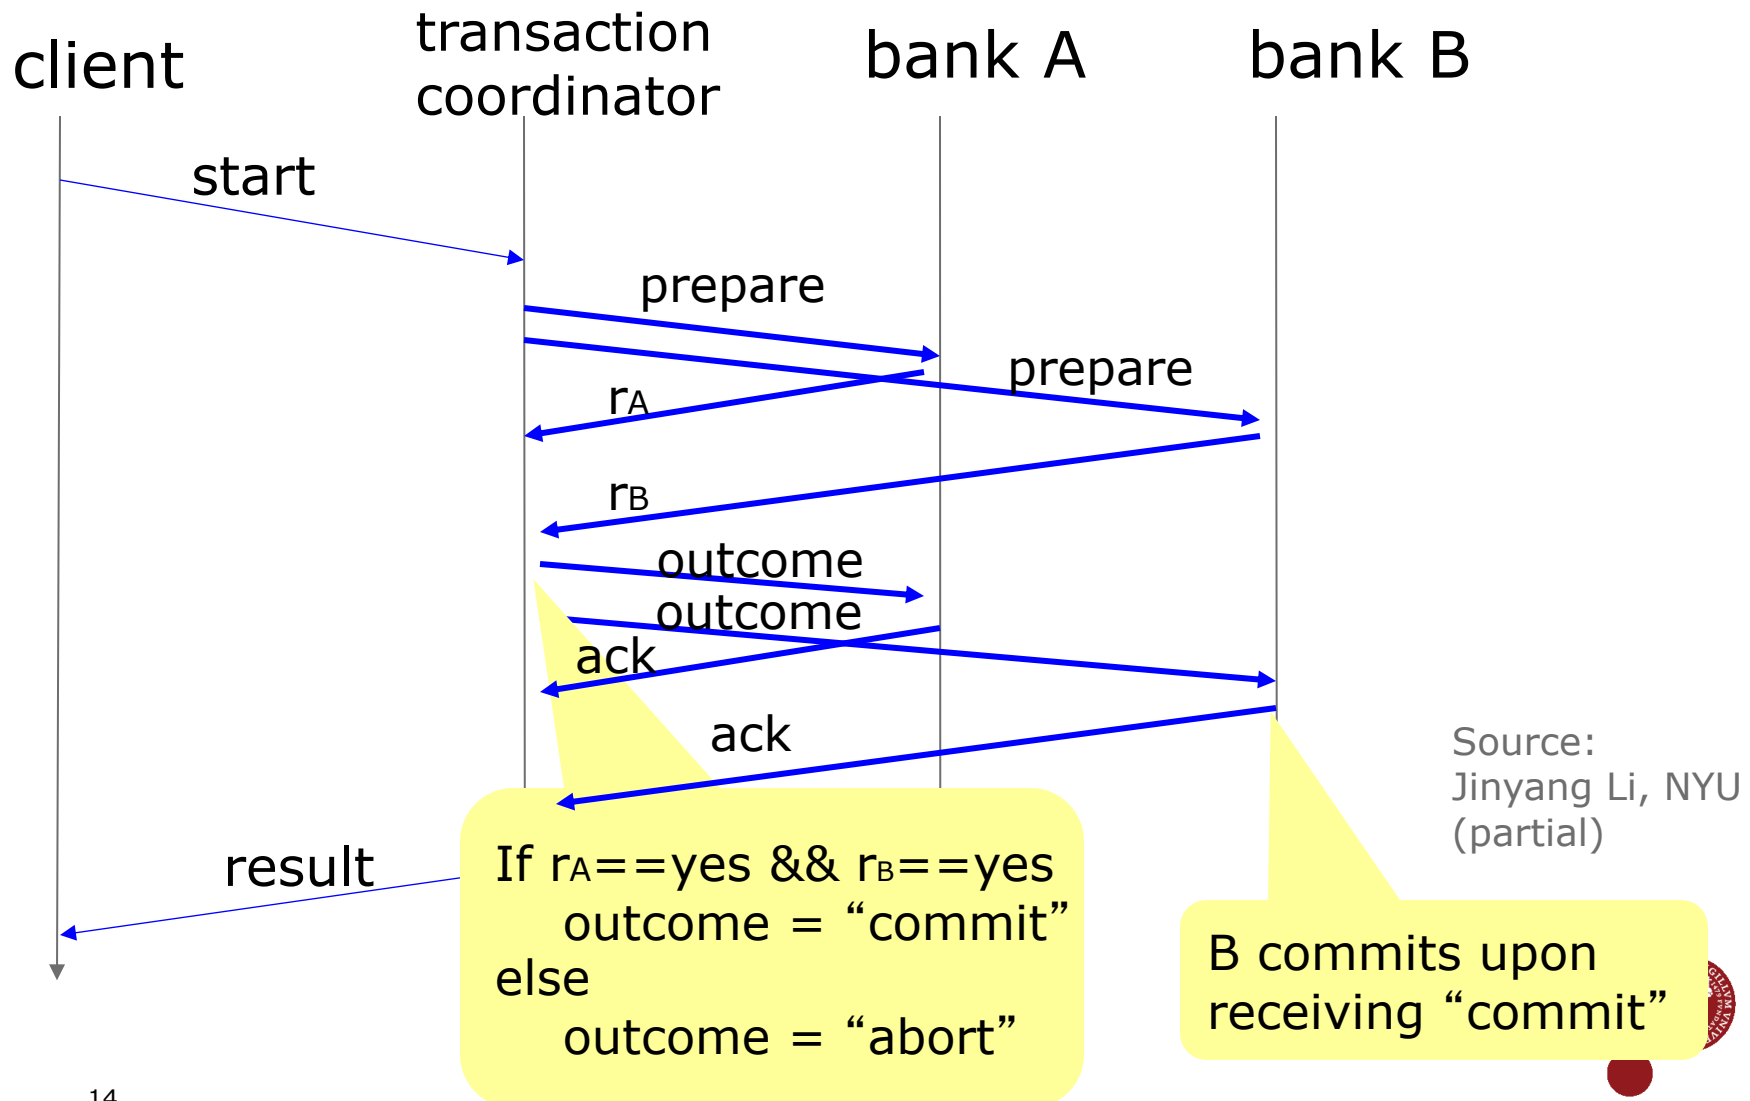
Transaction T1: TRANSFER

```
BEGIN
  UPDATE account
  SET bal = bal + 100
  WHERE account_id = 'A';
  --
  UPDATE account
  SET bal = bal - 100
  WHERE account_id = 'B';
COMMIT
```

- What if we partition accounts A and B across different computers?
- How would you execute the TRANSFER transaction using Message-Oriented Middleware (MOM)?
- How to achieve ACID?




Recall: Two-Phase Commit (2PC)



The CAP Theorem

- **Consistency:** client perceives set of operations occurred all at once (i.e., atomicity)
- **Availability:** every request must result in an intended response
- **Partition tolerance:** operations terminate, even if the network is partitioned.



CAP Theorem:
A scalable service cannot achieve all three properties simultaneously

ACID

- Using 2PC guarantees atomicity (i.e. C in CAP theorem)
- If 2PC is used, a transaction is not guaranteed to complete in the case of network partition (either Abort or Blocked)
 - i.e. we choose C over A
- But what if an application can benefit from choosing A over C
 - How to achieve this?



From ACID to BASE

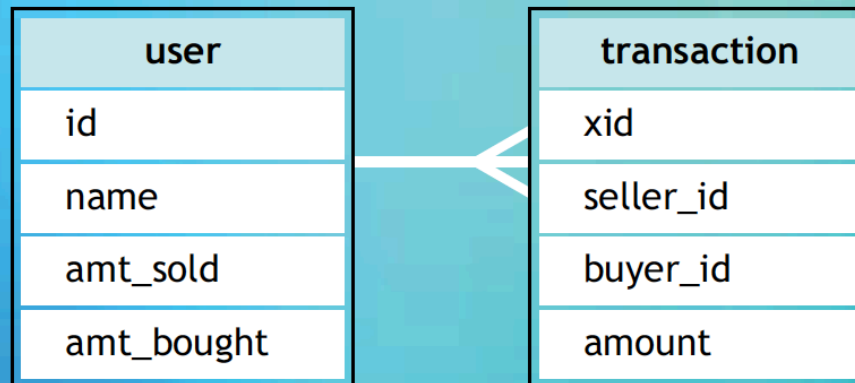
- **BASE**
 - **Basically-Available**: only components affected by failure become unavailable, not whole system
 - **Soft-State**: a component's state may be out-of-date, and events may be lost without affecting availability
 - **Eventually Consistent**: under no further updates and no failures, partitions converge to consistent state



A BASE Scenario

- Users buy and sell items
- Simple transaction for item exchange:

Sample Schema



Begin transaction


Insert into transaction(xid, seller_id, buyer_id, amount);

Update user set amt_sold=amt_sold+\$amount where id=\$seller_id;

Update user set amt_bought=amount_bought+\$amount where id=\$buyer_id;

End transaction

Decouple Item Exchange with Queues



Record Transaction
Queue User updates

Peek Queue for message
Check if update applied
If not

Update User amounts

Record update as applied

Remove message from queue

- **Issues**

- Tolerance to loss
- Idempotence
- Order



More Types of Communication Abstractions

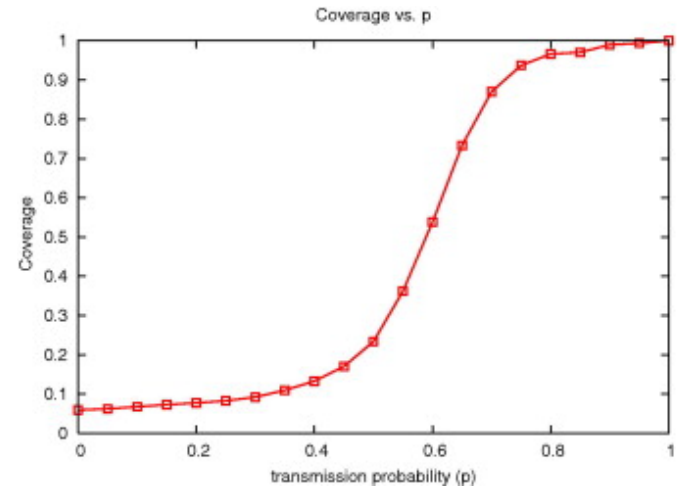
- **Stream-Oriented**
 - Continuous vs. discrete
 - Asynchronous vs. Synchronous vs. Isochronous
 - Simple vs. complex
- **Multicast**
 - SEND/RECEIVE over groups
 - Application-level multicast vs. gossip

More details in
compendium




A Word about Gossip

- **Epidemic protocols**
 - No central coordinator
- **Anti-entropy**
 - Each node communicates with random node
 - Round: every node does the above
 - Pull vs. push vs. both
 - Spreading update from single to all nodes takes $O(\log(N))$
- **Gossiping**
 - If P is updated it tells a random Q
 - If Q already knows, P can lose interest with some percentage
 - Not guaranteed that all nodes get infected by update



Coverage vs. trans prob




Efficiency vs. trans prob



Source: Haas et al. (partial)

What should we learn today?



- Describe different approaches to design communication abstractions, e.g., transient vs. persistent and synchronous vs. asynchronous
- Explain the design and implementation of message-oriented middleware (MOM)
- Explain how to organize systems employing a BASE methodology
- Discuss the relationship of BASE to eventual consistency and the CAP theorem
- Identify alternative communication abstractions such as data streams and multicast / gossip

Merry Christmas and a Happy New Year! 😊

