

# Welcome to ACS TA session 4

Julian, Leonardo, Li, Nikolaus, Rodrigo, Tilman, Yijian

Department of computer science, University of Copenhagen

Academic year 2021-2022, Block 2



# Agenda for today

## ❖ Revisit ARIES

- ARIES properties, approach, principles, data structures, phases
- Exercises on the ARIES algorithm

# Revisit ARIES: principles and properties

## Properties

- **Atomicity**: undo the transactions that do not commit
- **Durability**: ensure all actions of committed transactions survive system crashes and media failures

## Approach

- **Steal**: pages are written to disk in yet uncommitted transactions
- **No-force**: when a transaction commits, its pages are not forced to disk

## Principles

- **Write-ahead logging**: log the operation before executing it
- **Repeat history**: re-bring system to its state when it crashed and then fix
- **Log the undo**: to fully repeat the history, including the undo operations.

BUT we never undo the undo operations.

# ARIES log record data structure

## Log record

- **Log**: chronologic sequence of log entries
- **Log tail**: the portion of the log in main memory (not forced yet)
- **Log sequence numbers** (LSN): strictly increasing IDs for log records

## Log record types and fields

- *All*: *prevLSN*, *transID*, *type*
- **Update**: *pageID*, *length*, *offset*, *before-image*, *after-image*
- **Commit**, **Abort**, **End**
- **Compensation** (CLR): *undoNextLSN*

# ARIES additional data structures

## Dirty pages table

- One entry per page not written to disk yet
- **Fields:** pageID, recLSN

## Transactions table

- One entry per transaction
- **Fields:** transID, lastLSN, status
- Entries with status **committed** or **aborted** are removed from the

Table when the corresponding transaction reaches the **end** state

# ARIES phases

## Analysis

- Identifies dirty pages and active transactions

## Redo

- Repeats all actions from safe point to moment of crash
- Leaves the data structures in the latest state prior to the crash

## Undo

- Undoes the actions of uncommitted transactions reverse-chronologically

## Exercise 1

- 1) If we can guarantee that uncommitted data is never written to disk, is **undo** still necessary?
- 2) What about **redo**?
- 3) If updates are always forced to disk when a transaction commits, is **undo** still necessary?
- 4) What about **redo**?

## Exercise 1

1) If we can guarantee that uncommitted data is never written to disk, is **undo** still necessary?

2) What about **redo**? **No Steal** – **No Undo**

3) If updates are always forced to disk when a transaction commits, is **undo** still necessary?

4) What about **redo**?



## Exercise 1

1) If we can guarantee that uncommitted data is never written to disk, is **undo** still necessary?

2) What about **redo**? **No Steal** – **No Undo**

3) If updates are always forced to disk when a transaction commits, is **undo** still necessary?

4) What about **redo**? **Force** – **No Redo**

## ARIES Questions

After a crash failure, where in the log ...

1. should the **analysis** phase start?
2. should the **redo** phase start?
3. should the **undo** phase start?

## ARIES Questions

After a crash failure, where in the log ...

1. should the **analysis** phase start?

most recent checkpoint

2. should the **redo** phase start?

3. should the **undo** phase start?

## ARIES Questions

After a crash failure, where in the log ...

1. should the **analysis** phase start?

most recent checkpoint

2. should the **redo** phase start?

smallest recLSN in dirty page table

3. should the **undo** phase start?

# ARIES Questions

After a crash failure, where in the log ...

1. should the **analysis** phase start?

most recent checkpoint

2. should the **redo** phase start?

smallest recLSN in dirty page table

3. should the **undo** phase start?

largest lastLSN in transaction table

## Exercise 2

Apply the ARIES recovery algorithm to the next scenario. Show:

1. the state of the transaction and dirty page tables after the **analysis** phase
2. the sets of winner and loser transactions
3. the values for the LSNs where **redo** starts and **undo** ends

How far back into the log must ARIES scan during **redo** and **undo**?

What are:

4. the set of log records that may cause pages to be rewritten during redo?
5. the set of log records undone during undo?
6. the contents of the log after the recovery procedure completes?

# Exercise 2

**Xact  
table**

transID	status	lastLSN
T2	running	2
T1	committed	3

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4 5			begin checkpoint end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				

## Exercise 2

1. the state of the transaction and dirty page tables after the **analysis** phase

**Xact  
table**

transID	status	lastLSN
T2	running	2
T1	committed	3

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4 5			begin checkpoint end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				



## Exercise 2

1. the state of the transaction and dirty page tables after the **analysis** phase

**Xact  
table**

transID	status	lastLSN
T3	running	7

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2
A	7

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4 5			begin checkpoint end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				

# Exercise 2

## 2.the sets of winner and loser transactions

**Xact  
table**

transID	status	lastLSN
T3	running	7

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2
A	7

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4 5			begin checkpoint end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				

# Exercise 2

## 2.the sets of winner and loser transactions

**Xact  
table**

transID	status	lastLSN
T3	running	7

Winner: T1, T2

Loser: T3

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2
A	7

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4 5			begin checkpoint end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				

## Exercise 2

3. the values for the LSNs where **redo** starts and **undo** ends

**Xact  
table**

transID	status	lastLSN
T3	running	7

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2
A	7

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4			begin checkpoint	
5			end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				

## Exercise 2

3. the values for the LSNs where **redo** starts and **undo** ends

**Xact  
table**

transID	status	lastLSN
T3	running	7

**Redo start**

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2
A	7

**Undo ends**

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4			begin checkpoint	
5			end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				

## Exercise 2

4. the set of log records that may cause pages to be rewritten during redo?

**Xact  
table**

transID	status	lastLSN
T3	running	7

**Redo start**

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2
A	7

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4 5			begin checkpoint end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				

## Exercise 2

4. the set of log records that may cause pages to be rewritten during redo?

Xact  
table

transID	status	lastLSN
T3	running	7

Redo start

Dirty  
page  
table

pageID	recLSN
C	1
B	2
A	7

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4 5			begin checkpoint end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
CRASH!!!				

## Exercise 2

5.the set of log records undone during undo?

**Xact  
table**

transID	status	lastLSN
T3	running	7

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2
A	7

**Undo ends**

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4 5			begin checkpoint end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				



## Exercise 2

5.the set of log records undone during undo?

**Xact  
table**

transID	status	lastLSN
T3	running	7

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2
A	7

**Undo ends**

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4			begin checkpoint	
5			end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				

## Exercise 2

6. the contents of the log after the recovery procedure completes?

**Xact  
table**

transID	status	lastLSN
T3	running	7

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2
A	7

ToUndo:

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4 5			begin checkpoint end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				

## Recovery: The UNDO Phase

ToUndo = { lsn | lsn a lastLSN of a “loser” Xact }

### Repeat:

- Choose largest LSN among ToUndo.
- If this LSN is a CLR and undonextLSN == NULL
  - Write an End record for this Xact.
- If this LSN is a CLR, and undonextLSN != NULL
  - Add undonextLSN to ToUndo
- Else this LSN is an update. Undo the update, write a CLR, add prevLSN to ToUndo.

Until ToUndo is empty.



6.the contents of the log after the recovery procedure completes?

## Exercise 2

**Xact  
table**

transID	status	lastLSN
T3	running	7

**Dirty  
page  
table**

pageID	recLSN
C	1
B	2
A	7

ToUndo: 7

LSN	prevLSN	transID	type	pageID
1	null	T1	update	C
2	null	T2	update	B
3	1	T1	commit	
4			begin checkpoint	
5			end checkpoint	
6	3	T1	end	
7	null	T3	update	A
8	2	T2	update	C
9	8	T2	commit	
10	9	T2	end	
<b>CRASH!!!</b>				
11	7	T3	abort	
12	11	T3	CLR:Undo LSN 7	
13	12	T3	end	

## Exercise 3

The next scenario depicts a situation where the system crashes during recovery. Apply the ARIES algorithm after:

1. **Crash 1** that occurred during normal execution
2. **Crash 2** that occurred during recovery

## Exercise 3

LSN	prevLSN	transID	Type	undoNextLSN	pageID
01 05			Begin checkpoint End checkpoint		
10		T1	Update		P5
20		T2	Update		P3
30	10	T1	Abort		
40 45	30 40	T1	CLR:Undo LSN 10 End		
50		T3	Update		P1
60	20	T2	Update		P5
<b>CRASH 1 !!!</b>					
70	60	T2	Abort		
80	50	T3	Abort		
90	70	T2	CLR:Undo LSN 60	20	
100 105	80 100	T3	CLR:Undo LSN 50 End		
<b>CRASH 2 !!!</b>					

# Thank you!

Julian, Leonardo, Li, Nikolaus, Rodrigo, Tilman, Yijian

Department of computer science, University of Copenhagen

Academic year 2021-2022, Block 2

